

Climate in the cloud: Democratizing access to climate projections

OAR Cloud Tiger Webinar Series
V.Balaji, Aparna Radhakrishnan



PRINCETON
UNIVERSITY



Sharing our (ongoing) journey to the cloud.

- Setting up cloud storage
- Accessing cloud storage
 - Preparation for ESGF data publication in the cloud
 - Resources to foster analysis in the cloud
- Monitoring and cost exploration
- Acknowledgments



Cloud P.C. princeton.edu

P.C. Adih Haarish

Setting up Cloud Storage

- ❑ Storage service used: Amazon Simple Storage Service (Amazon S3)
- ❑ Automatic cost savings with Intelligent-tiering enabled
 - Optimized for frequent access
 - Lower cost tier optimized for infrequent access
- ❑ Security policies: S3 bucket policies, object access and bucket access control lists (ACLs), IAM (Identity and Access Management) roles.
- ❑ Objects organized compliant to CMIP6 Data Reference Syntax (DRS)

e.g.

/CMIP6/CMIP/NOAA-GFDL/GFDL-CM4/1pctCO2/r1i1p1f1/Amon/ct/gr1/v20180701/



CMIP6 NetCDF model output from the ESGF public archives

Ack: CMIP6 data producers and coordinators worldwide

Accessing cloud storage (for development)

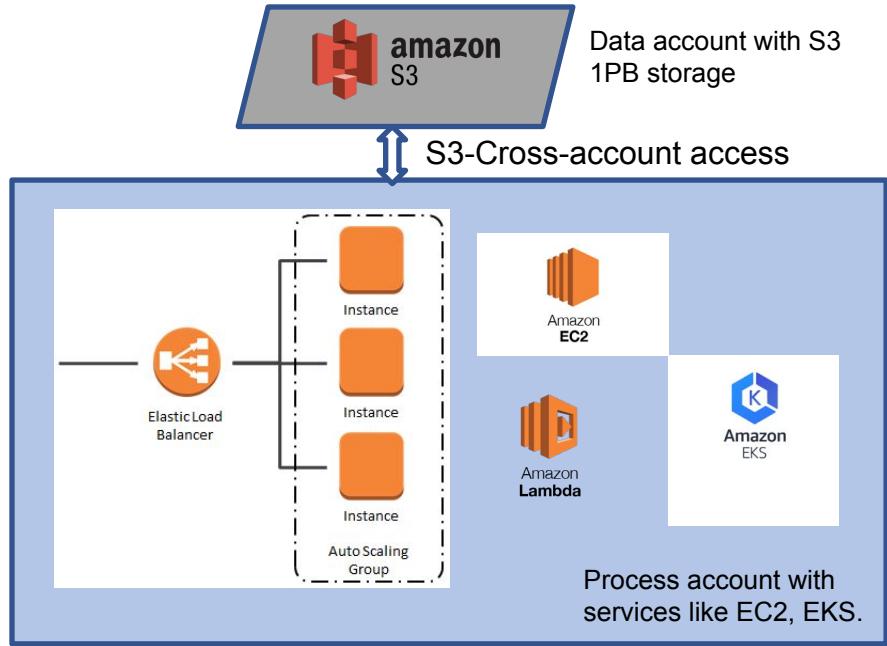
- We have 2 AWS awards/accounts from the ASDI initiative.

- Data account
- Process account

We need to enable cross-account access for S3 operations.

Few available options:

1. AWS automatic replication to the new account
2. Role switching: allows you to assume the role and permissions of another account
3. Cross-account permission with resource based policies and IAM policies



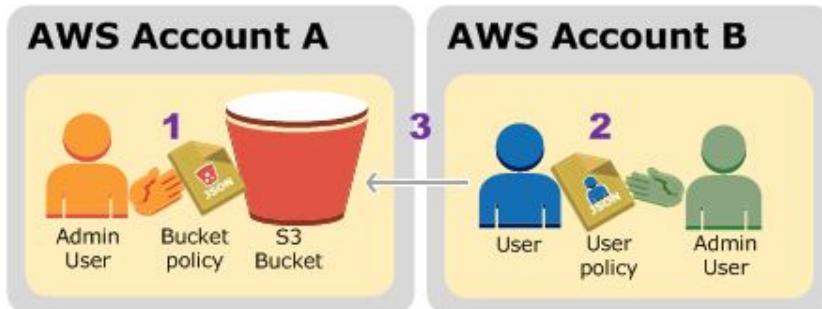
Cross Account Permissions

- We adopted the use of granting cross account permissions updating bucket policies and IAM role policy.
- This gives authorized users in the second account the credentials to perform select S3 operations in the data account with the same permissions as the original account

Example: aws s3 ls s3://esgf-world/CMIP6/ --profile processaccount

```
PRE C4MIP/  
PRE CDRMIP/  
PRE CFMIP/  
PRE CMIP/  
PRE DAMIP/  
PRE FAFMIP/  
PRE GMMIP/  
PRE HighResMIP/  
PRE LUMIP/  
BBE OMTD/
```

Example that uses **AWS CLI commands** to list, copy data to the data account from the process account profile. AWS CLI was used for data transfer operations in the first phase of the project



Ref

<https://docs.aws.amazon.com/AmazonS3/latest/dev/example-walkthroughs-managing-access-example2.html>

Accessing cloud storage (for development)

Use-case: The existing tools such as the **ESGF software stack Data publication project** in AWS requires interaction with the Amazon S3 bucket to perform R/W operation

Our team explored options to mount S3 as a local drive
s3fs

A FUSE filesystem application backed by AWS.

pros: easy setup

cons: Latency when performing IO centric operations and metadata operations such as listing/searching. TBD:
cost-efficiency.

goofys:

A high-performance, POSIX-ish Amazon S3 file system written in Go

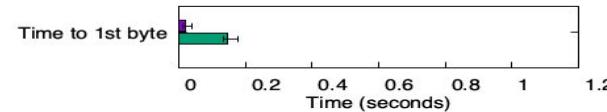
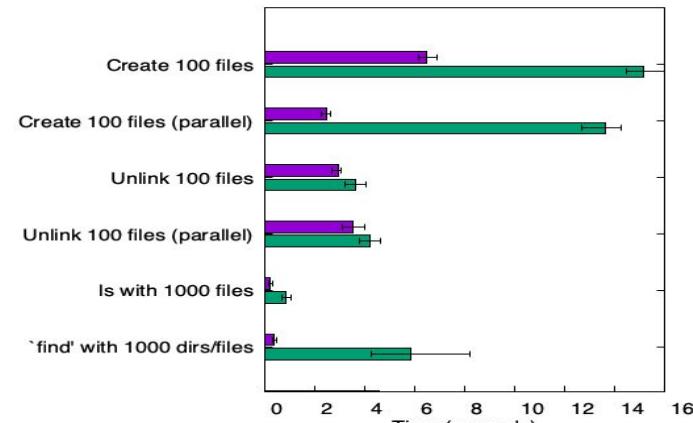
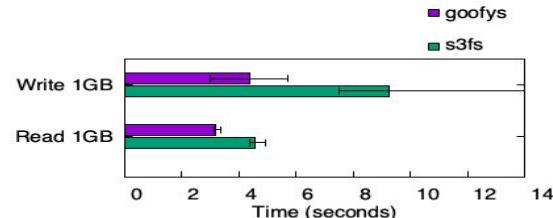
pros: easy setup, in-memory cache

cons: TBD

Amazon Lustre FSx: Amazon High-performance file system.

pros: High Performance

cons: Redundant copy of S3, NOT cost efficient for our use-case



Benchmarking results
from
<https://github.com/kaing/goofys>

Building blocks for analysis: The GIST



photo credit: HV

- Gentle Learning curve
- Interoperability
- Scalability Shareability
- Traceability



Community-driven development
collaborations



Climate in the cloud, V.Balaji, A.Radhakrishnan

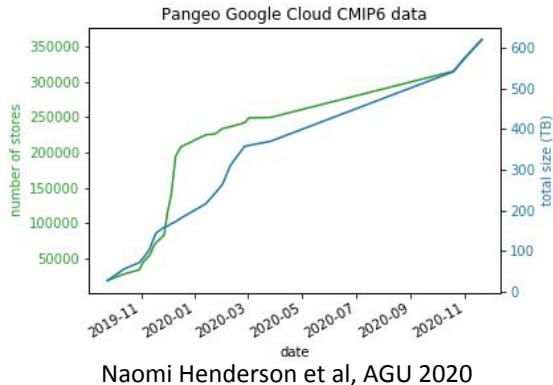


Community-driven development

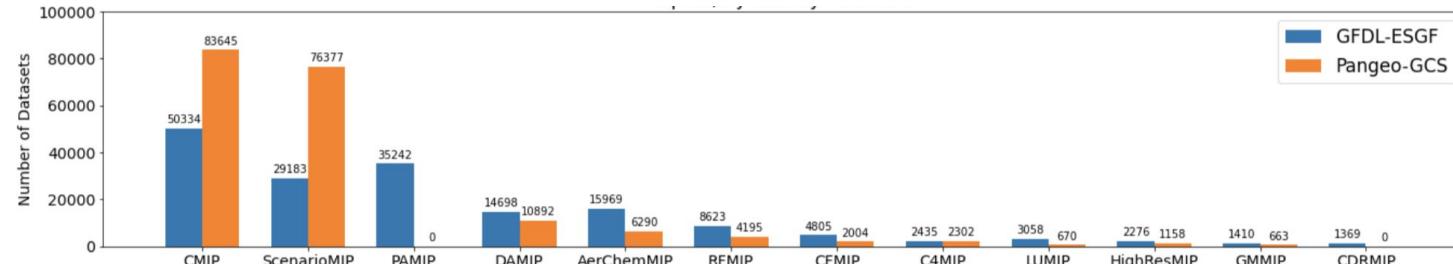
A **community of people** working collaboratively to develop **software and infrastructure** to enable **Big Data geoscience research**.



Mission



To **cultivate an ecosystem** in which the next generation of open-source analysis tools for ocean, atmosphere and climate science can be developed, distributed, and sustained. These tools must be **scalable** in order to meet the current and future challenges of big data, and these solutions should leverage the existing expertise outside of the geoscience community.



Ack: Ryan Abernathey and Pangeo team.
<https://pangeo.io/about.html#about-pangeo>

Fig courtesy: Naomi Henderson et al, October 2020.



Climate in the cloud, V.Balaji, A.Radhakrishnan



Xarray

- A high-level API for loading, transforming, and performing calculations on multi-dimensional arrays.
- Built leveraging NumPy and pandas API
- Code it like you say it (Easy to get started!)
 - E.g. `ds.sel(time='2000-01') ds['thetao'].sel(z_l=2.5).mean(dim='time')`
- Incentive, motivation to adhere to CF conventions.
 - Leverages the use of **CF metadata** Conventions
- Simple gateway to exploring **different data formats and input sources**.
 - E.g. NetCDF, OPeNDAP, Google cloud data store, Zarr,.....
- xarray's data structures can be backed by **dask**

*xarray: originally developed by
S. Hoyer.*

<https://github.com/pydata/xarray>



Climate in the cloud, V.Balaji, A.Radhakrishnan



Gentle learning curve



DASK



- Fits into a cohesive ecosystem
- Multi-DASKS: Scales up (distributed clusters) and down (single-machine scheduler)
- Less coding intervention to make (many) script Dask-able.
- Provides effective triaging and monitoring system (Dask dashboard),..

Dask, originally developed by Matthew Rocklin
<https://dask.org/>



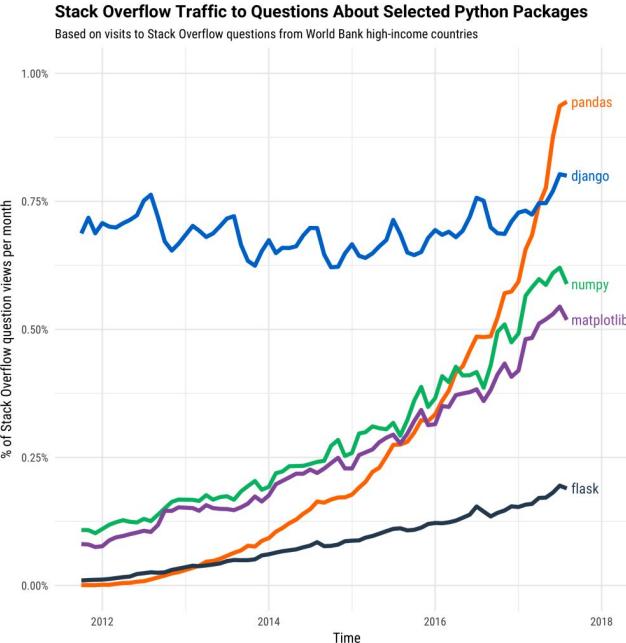
Scalable
Inter-operable
Gentle learning curve



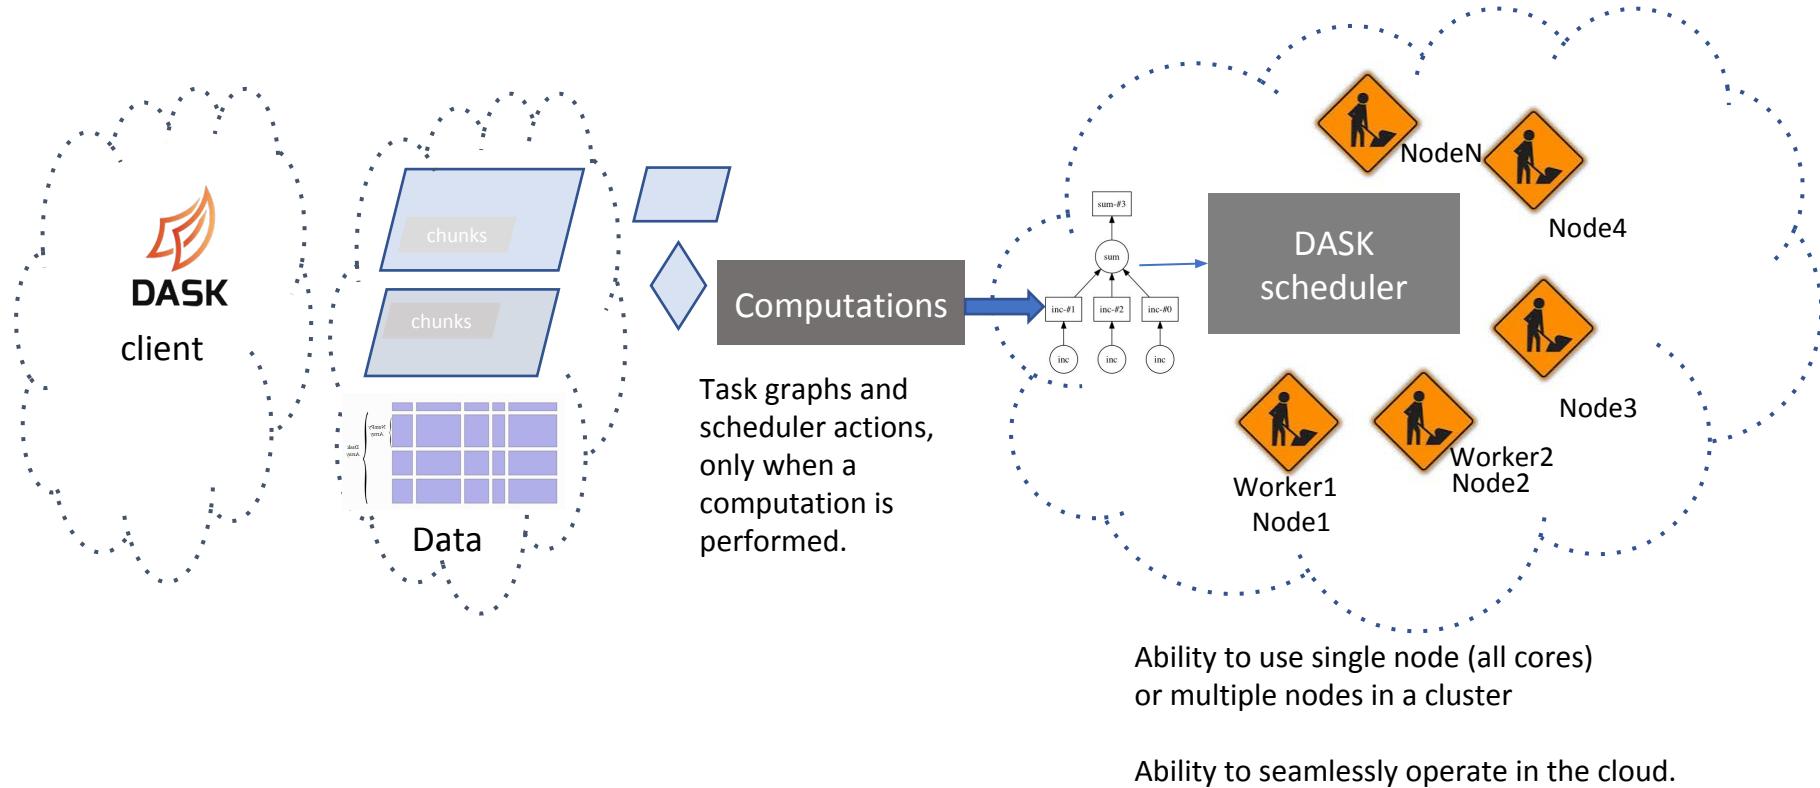
Climate in the cloud, V.Balaji, A.Radhakrishnan



Image credit to Stack Overflow blogposts and R. Abernathy



Chain of actions in (lazy) Dask: A bird's-eye view



Zarr



Zarr

- A storage format optimized for high throughput distributed reads on multi-dimensional arrays.
- Zarr works well on both traditional filesystem storage and on Cloud Object Storage.

Pluggable storage

`zarr.DirectoryStore`, `zarr.ZipStore`,
`zarr.DBMStore`, `zarr.LMDBStore`, `zarr.SQLiteStore`,
`zarr.MongoDBStore`, `zarr.RedisStore`,
`zarr.ABSStore`, `s3fs.S3Map`, `gcsfs.GCSMap`, ...

Zarr Scalable Storage of Tensor Data for Use in Parallel and Distributed Computing, SciPy 2019, A. Miles

Buckets / example_rdussin / OM4p5_sample_store / uo

<input type="checkbox"/>	<input type="checkbox"/> .zarray	335 B	application/octet-stream	Standard
<input type="checkbox"/>	<input type="checkbox"/> .zattrs	393 B	application/octet-stream	Standard
<input type="checkbox"/>	<input type="checkbox"/> 0.0.0.0	25.81 MB	application/octet-stream	Standard
<input type="checkbox"/>	<input type="checkbox"/> 1.0.0.0	25.81 MB	application/octet-stream	Standard

<https://github.com/raphaeldussin/MOM6-AnalysisCookbook>



Climate in the cloud, V.Balaji, A.Radhakrishnan



Examples: Dask and xarray

Use dask.distributed task scheduler and launch DASK using SLURMcluster

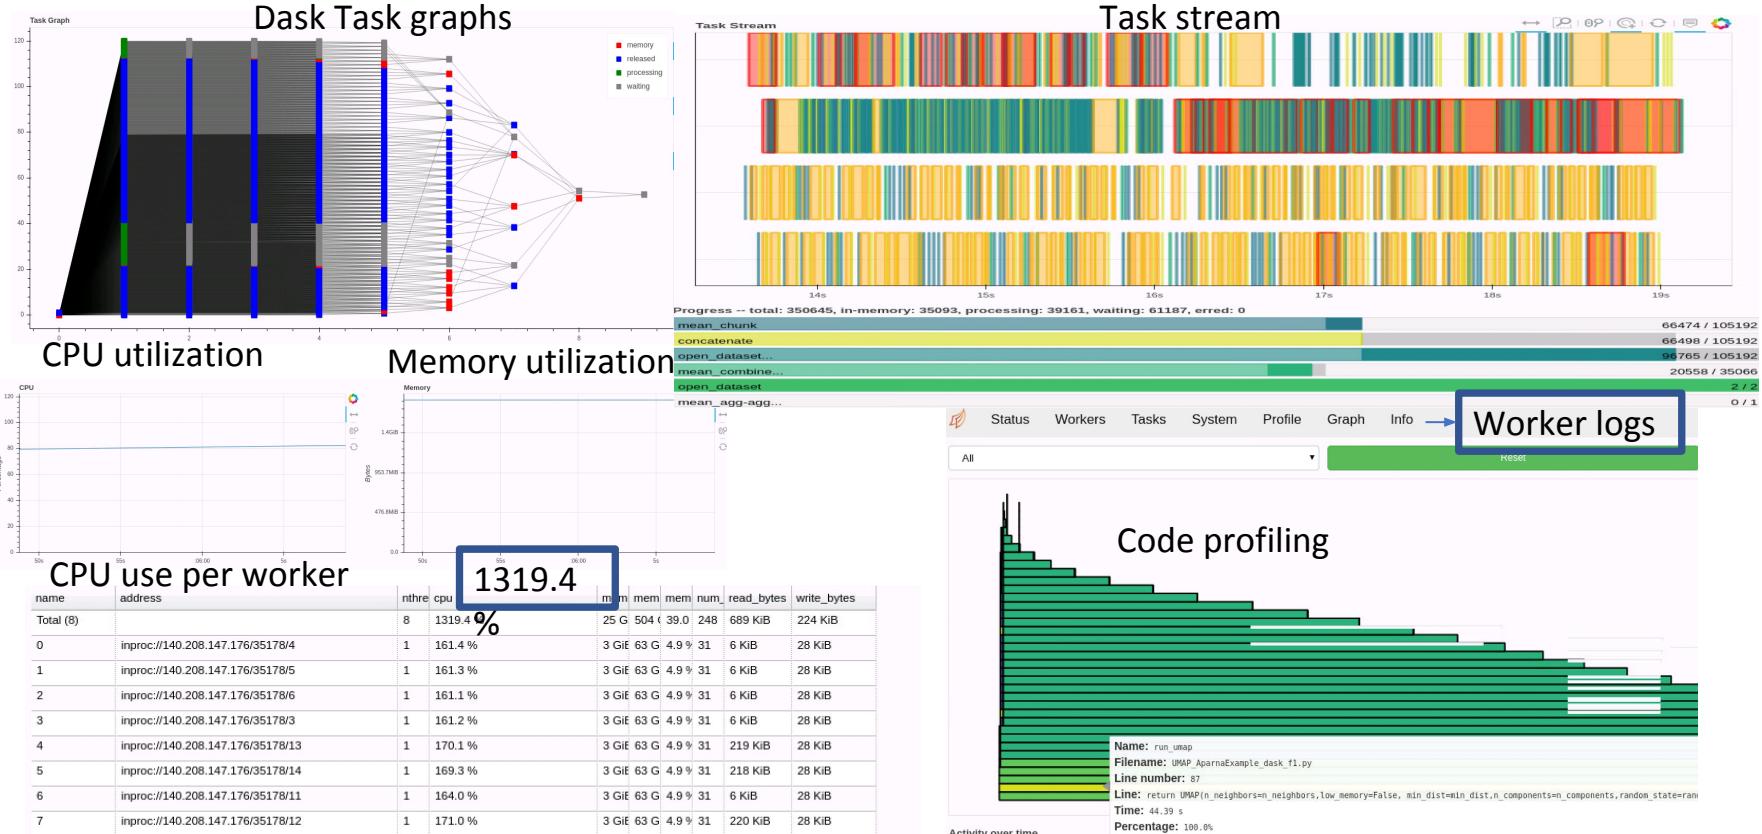
```
[3]: from dask.distributed import Client

#Instantiate Dask client
if (clusterType == "local"):
    from dask.distributed import LocalCluster
    cluster = LocalCluster(dashboard_address=dashPort,local_directory=localdir)
else:
    from dask_jobqueue import SLURMCluster
    scheduler_options = {}
    scheduler_options["dashboard_address"] = dashPort
    cluster = SLURMCluster(queue='batch',memory=mem,project='gfdl_f',cores=numCores,walltime='2:60:00',
                           scheduler_options=scheduler_options,log_directory=logdir,
                           local_directory=(os.getenv('TMPDIR')))

cluster.scale(numWorkers)
client = Client(cluster)
client
```



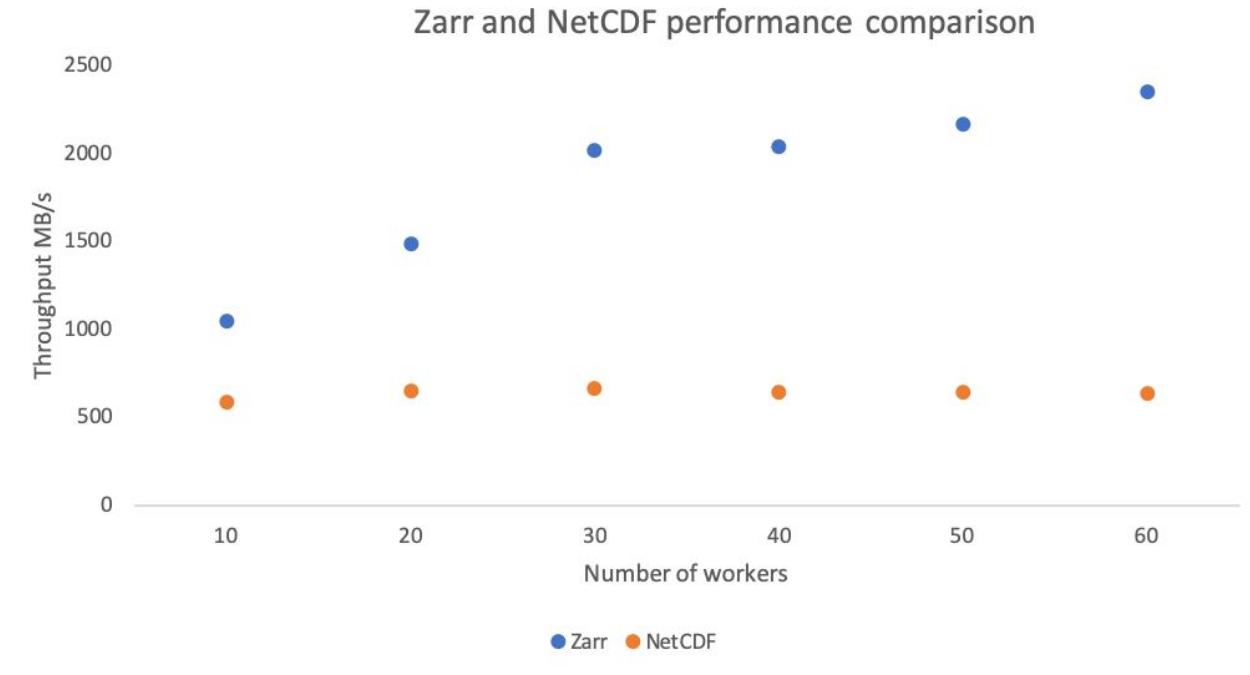
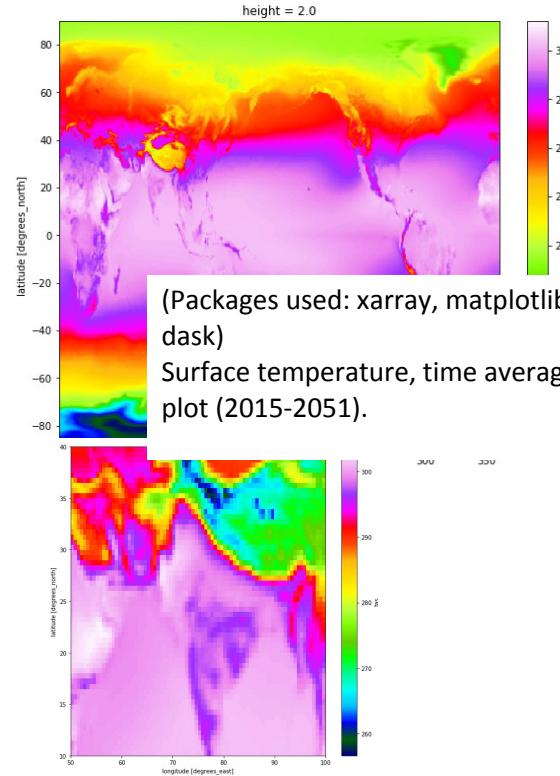
Dask dashboard



Climate in the cloud, V.Balaji, A.Radhakrishnan



Preliminary Benchmarking Results



Ack: Zhao, Ming; et al.

<https://doi.org/10.22033/ESGF/CMIP6.2262>

2015-2051



Climate in the cloud, V.Balaji, A.Radhakrishnan

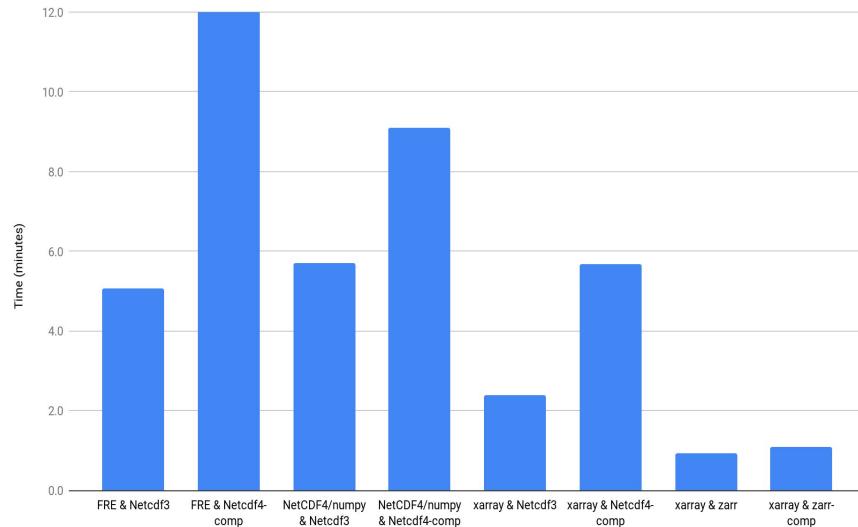
Ref. <https://github.com/aradhakrishnanGFDL/enes2020>



Preliminary Benchmarking Results (contd..)

2D tos SST notebook timing varying method & data format

Create and plot 300-year climatology from two Op25 and three Op5 simulations on an101

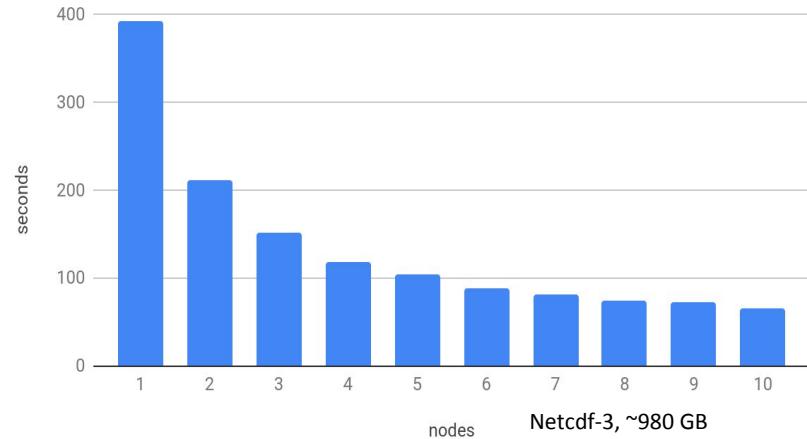


zarr was faster than NetCDF

compressed zarr was *much* faster than compressed NetCDF

Resource scaling test

Create and plot 300-year climatology from Op25 and Op5 simulations on PP/AN cluster



Perry et al, 2019



Examples: Dask and scikit-learn

Before DASK

```
: X, y = make_blobs(n_samples = 150000, n_features = 2, centers = 2, cluster_std = 1.9)
model = DBSCAN(eps = 0.5, min_samples = 20)
%time model.fit(X)
```

```
CPU times: user 6 s, sys: 638 ms, total: 6.64 s
Wall time: 6.61 s
```

AFTER DASK

```
cluster.adapt(minimum=1,maximum=4)
```

```
##Test with DASK
```

```
: from joblib import parallel_backend,parallel
X, y = make_blobs(n_samples = 150000, n_features = 2, centers = 2, cluster_std = 1.9)
model = DBSCAN(eps = 0.5, min_samples = 20,n_jobs=-1)
with parallel_backend('dask'):
    %time model.fit(X)
```

```
CPU times: user 8.73 s, sys: 563 ms, total: 9.29 s
Wall time: 3.92 s
```



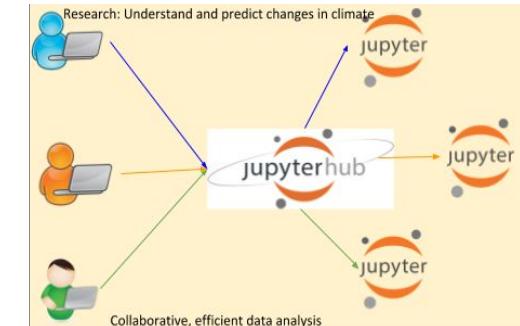
Climate in the cloud, V.Balaji, A.Radhakrishnan



JupyterHub



- The **Jupyter Notebook** is an open-source web application that allows you to create and share documents that contain live code and documentations.
- JupyterHub** brings the power of Jupyter notebooks to groups of users.
- JupyterHub gives users access to computational environments and resources without burdening the users with installation and maintenance tasks.
- Ideal for “Bringing desktop to data”



The screenshot shows the JupyterHub interface running on AWS. On the left is a file browser with a sidebar containing icons for file operations like upload, download, and refresh. The main area displays a list of Jupyter notebooks:

- gfdl-aws-analysis / examples /
- Name
- dask_netcdf_LittlestJupyter.ipynb
- intake-esm-s3-nc-simple.ipynb
- intake-esm-s3-nc-test.ipynb
- s3_list_example.py
- S3-public-access-usage-...
- successful_script.ipynb
- Untitled.ipynb
- xarray-dask-testing-on-E...

Annotations point to specific parts of the interface:

- An arrow points to the explanatory text and annotations in the notebook cell below.
- An arrow points to the code in the notebook cell below.

Below the file browser, the toolbar includes icons for file operations like new, open, save, and refresh, along with a dropdown menu for Markdown and a tab labeled "coolenv".

The notebook cell contains the following content:

```
[1]: from netCDF4 import Dataset
import numpy as np
import matplotlib
from matplotlib import pyplot as plt
import xarray as xr
import intake,yaml
import intake_esm
import numpy as np
%matplotlib inline
```

Screenshot of JupyterHub deployed in our AWS

Ref. <https://jupyter.org/hub>

Amazon Elastic Compute Cloud - EC2



Amazon
EC2

provides scalable compute capacity on
the cloud

One EC2 instance can give you the
power to launch a simple JupyterHub

With the combined power of many EC2
instances you can speed workflow

The screenshot shows the AWS Lambda console interface. On the left, there's a sidebar with a 'Compute' icon and a list of services: Lightsail (with a hand cursor icon over it), Elastic Container Service, EKS, Lambda, Batch, and Elastic Beanstalk. To the right, there's a large rectangular area containing the text 'the littlest jupyterhub' with an orange logo, and a blue link below it: <https://tljh.jupyter.org/en/latest/install/amazon.html>.

Currently selected: t2.micro (- ECUs, 1 vCPUs, 2.5 GHz, -, 1 GiB memory, EBS only)

EC2 instance naming convention: <family><generation>.size>

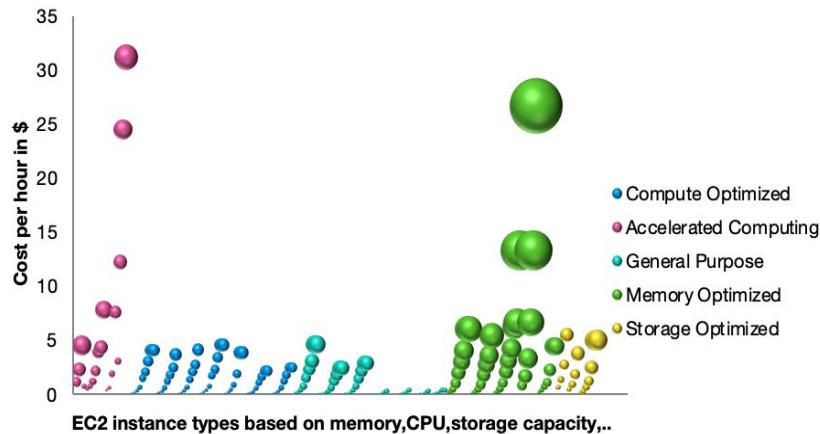
	Family	Type	vCPUs	Memory (GiB)	Instance Storage (GB)	Network Performance
<input type="checkbox"/>	t2	t2.nano	1	0.5	EBS only	Low to Moderate
<input checked="" type="checkbox"/>	t2	t2.micro <small>Free tier eligible</small>	1	1	EBS only	Low to Moderate
<input type="checkbox"/>	t2	t2.small	1	2	EBS only	Low to Moderate
<input type="checkbox"/>	t2	t2.medium	2	4	EBS only	Low to Moderate
<input type="checkbox"/>	t2	t2.large	2	8	EBS only	Low to Moderate
<input type="checkbox"/>	t2	t2.xlarge	4	16	EBS only	Moderate
<input type="checkbox"/>	t2	t2.2xlarge	8	32	EBS only	Moderate

Cost and usage of EC2

Choice of EC2 instances may depend on the **nature of your application** requirements, region specification and **cost** factors.

Scope for **optimization**: Auto-scaling, scheduled time for nodes.

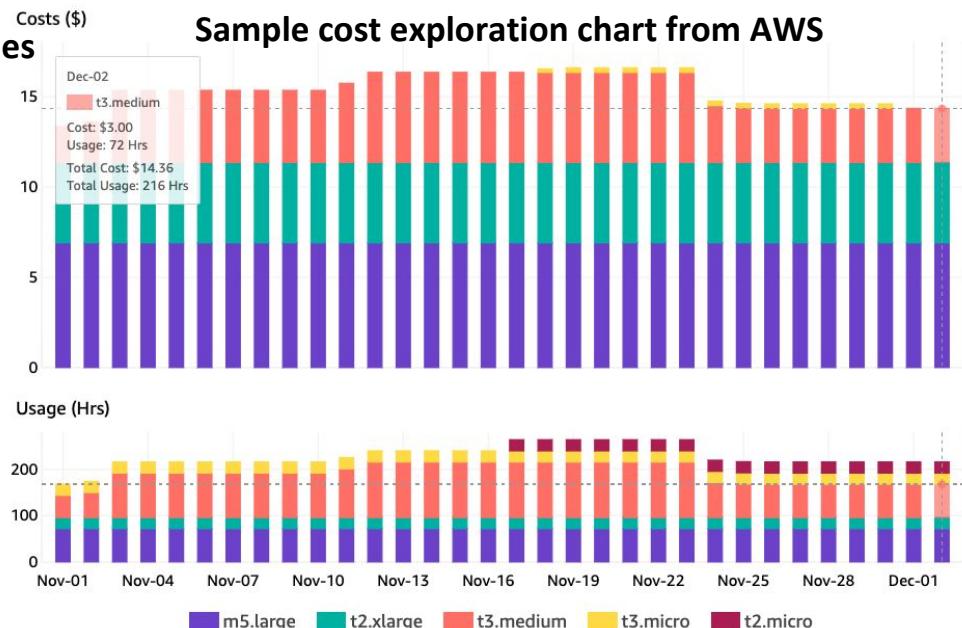
Cost comparison of different types of EC2 instance families



Disclaimer: Cost information on this slide may vary based on region, time, etc.

The values presented here are only approximate costs. For accurate info, please refer to your AWS account billing dashboard.

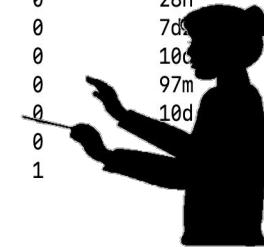
Sample cost exploration chart from AWS



Kubernetes and JupyterHub

- Kubernetes is a container orchestration system designed to facilitate containers and virtual machines for an application
- Kubernetes cluster in conjunction with JupyterHub can be configured to use multiple EC2 instances, which give a much larger compute power than our first attempt

```
(base) aos-pqlvdl:gfdl-aws-analysis ar46$ kubectl --namespace jhub get pods
NAME          READY   STATUS    RESTARTS   AGE
continuous-image-puller-58x82  1/1     Running   0          10d
continuous-image-puller-6vwg4   1/1     Running   0          10d
continuous-image-puller-8h6gm  1/1     Running   0          28h
continuous-image-puller-j4qhl  1/1     Running   0          10d
continuous-image-puller-m2r2x  1/1     Running   0          27h
continuous-image-puller-m9s65  1/1     Running   0          28h
continuous-image-puller-mvjft  1/1     Running   0          7d23h
continuous-image-puller-r97mw  1/1     Running   0          10d
continuous-image-puller-rwlb2  1/1     Running   0          28h
continuous-image-puller-xh9z1  1/1     Running   0          28h
continuous-image-puller-xvsq5  1/1     Running   0          7d
hub-dfbc895cf-xrp5z        1/1     Running   0          10d
jupyter-jupyter-5fadmin      1/1     Running   0          97m
proxy-75d76c4c64-bhs94       1/1     Running   0          10d
user-scheduler-5d97fdd64f-dxd6n 1/1     Running   0
user-scheduler-5d97fdd64f-sm2rf 1/1     Running   1
```



Amazon Elastic Kubernetes Service - EKS

Amazon has a built in way to run Kubernetes services called
EKS

- EKS can utilize multiple EC2 instances to create a cluster and provide compute power to multiple Kubernetes pods
- minikube (sets up Kubernetes cluster locally) could be employed for testing (we used it to test the ESGF software stack)
- EKS installation is based on this [tutorial](#) from AWS.
- Auto-scaling options can be enabled as needed.



Installing JupyterHub in EKS

Setup Kubernetes with required security permissions, e.g AWS EKS



Setup helm package manager



Install JupyterHub using [JupyterHub Helm Charts](#)



Update configurations to customize authentications, dask worker environments, etc



Run a helm upgrade to update

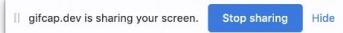
```
(base) aos-pqlvdl:gfdl-aws-analysis ar46$ helm list
```

NAME	REVISION	UPDATED	STATUS	CHART	APP VERSION	NAMESPACE
appt-release	1	Tue Aug 18 14:23:55 2020	DEPLOYED	dask-4.3.1	2.23.0	default
jhub	1	Sun Aug 9 00:43:22 2020	DEPLOYED	jupyterhub-0.9.0	1.1.0	jhub

JupyterHub in EKS



Sign in with GitHub



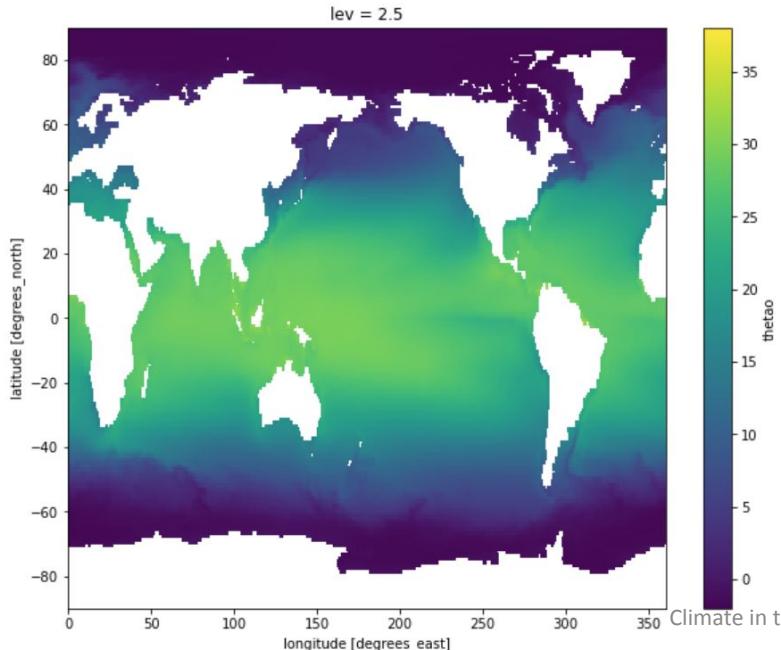
Performance comparison in the cloud/AWS

1 EC2 instance running dask local cluster in Littlest JupyterHub

```
CPU times: user 7min 5s, sys: 2min 41s, total: 9min 46s  
Wall time: 24min 49s
```

```
/opt/conda/lib/python3.7/site-packages/dask/array/numpy_compat.py:41: RuntimeWarning:  
divide by zero encountered in divide  
x = np.divide(x1, x2, out)
```

```
: <matplotlib.collections.QuadMesh at 0x7efff18f5cdd0>
```

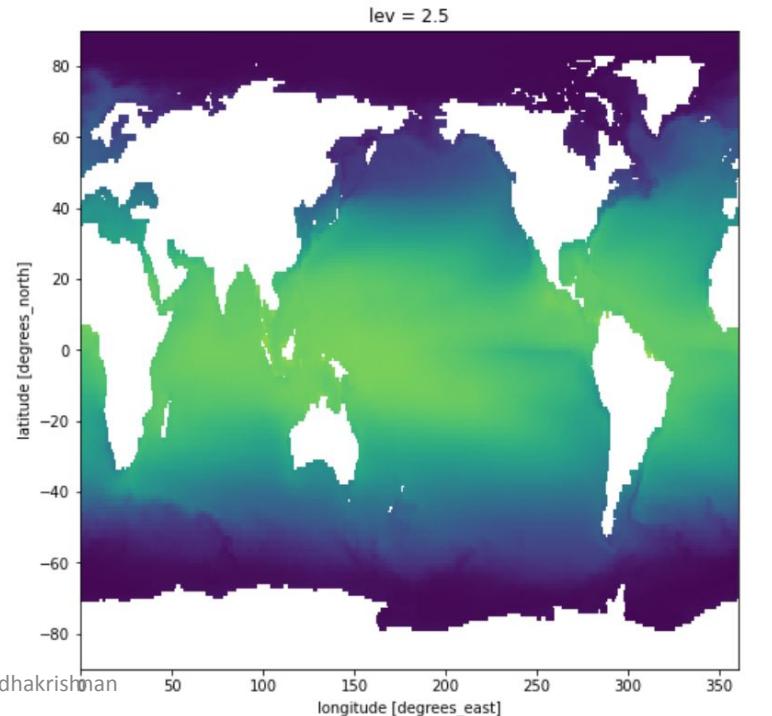


3 worker nodes/EC2 running Dask distributed on JupyterHub in EKS

Data citation: Adcroft, Alistair et al, (2018). NOAA-GFDL GFDL-CM4 model output prepared for CMIP6 OMIP. Version 20180701[1].Earth System Grid Federation. <https://doi.org/10.22033/ESGF/CMIP6.1403>

```
CPU times: user 1.77 s, sys: 53.2 ms, total: 1.82 s  
Wall time: 5min 46s
```

```
Out[7]: <matplotlib.collections.QuadMesh at 0x7fd5076a5390>
```



Useful resources to access S3 using Python

1. [Quick-start-guide](#) to use our example [Jupyter notebooks](#).
2. [Example](#) script that shows how to list contents of an S3 bucket using python3, boto3.
3. [Example](#) to show how to open a dataset in S3 using s3fs in Python.

```
import s3fs
import xarray as xr
fs_s3 = s3fs.S3FileSystem(anon=True)
s3path =
's3://<bucketName>/CMIP6/OMIP/NOAA-GFDL/GFDL-CM4/omi
p1/r1i1p1f1/Omon/thetao/gr/v20180701/thetao_Omon_GFDL-
CM4_omip1_r1i1p1f1_gr_170801-172712.nc'
remote_file_obj = fs_s3.open(s3path, mode='rb')

xarray_dataset_object = xr.open_dataset(remote_file_obj,
chunks={'time': 1})
```

xarray.Dataset																							
► Dimensions:		(bnds: 2, lat: 180, lev: 35, lon: 360, time: 3600)																					
▼ Coordinates:																							
<table><tr><td>lon</td><td>(lon)</td><td>float64</td><td>0.5 1.5 2.5 ... 357.5 358.5 359.5</td></tr><tr><td>lat</td><td>(lat)</td><td>float64</td><td>-89.5 -88.5 -87.5 ... 88.5 89.5</td></tr><tr><td>lev</td><td>(lev)</td><td>float64</td><td>2.5 10.0 20.0 ... 6e+03 6.5e+03</td></tr><tr><td>time</td><td>(time)</td><td>object</td><td>1708-01-16 12:00:00 ... 2007-12-...</td></tr></table>				lon	(lon)	float64	0.5 1.5 2.5 ... 357.5 358.5 359.5	lat	(lat)	float64	-89.5 -88.5 -87.5 ... 88.5 89.5	lev	(lev)	float64	2.5 10.0 20.0 ... 6e+03 6.5e+03	time	(time)	object	1708-01-16 12:00:00 ... 2007-12-...				
lon	(lon)	float64	0.5 1.5 2.5 ... 357.5 358.5 359.5																				
lat	(lat)	float64	-89.5 -88.5 -87.5 ... 88.5 89.5																				
lev	(lev)	float64	2.5 10.0 20.0 ... 6e+03 6.5e+03																				
time	(time)	object	1708-01-16 12:00:00 ... 2007-12-...																				
▼ Data variables:		<table><tr><td>lat_bnds</td><td>(time, lat, bnds)</td><td>float64</td><td>dask.array<chunksize=(240, 180, 2), meta=np.ndarray></td></tr><tr><td>lon_bnds</td><td>(time, lon, bnds)</td><td>float64</td><td>dask.array<chunksize=(240, 360, 2), meta=np.ndarray></td></tr><tr><td>thetao</td><td>(time, lev, lat, lon)</td><td>float32</td><td>dask.array<chunksize=(1, 35, 180, 360), meta=np.ndarray></td></tr><tr><td>time_bnds</td><td>(time, bnds)</td><td>object</td><td>dask.array<chunksize=(1, 2), meta=np.ndarray></td></tr><tr><td>lev_bnds</td><td>(time, lev, bnds)</td><td>float64</td><td>dask.array<chunksize=(240, 35, 2), meta=np.ndarray></td></tr></table>		lat_bnds	(time, lat, bnds)	float64	dask.array<chunksize=(240, 180, 2), meta=np.ndarray>	lon_bnds	(time, lon, bnds)	float64	dask.array<chunksize=(240, 360, 2), meta=np.ndarray>	thetao	(time, lev, lat, lon)	float32	dask.array<chunksize=(1, 35, 180, 360), meta=np.ndarray>	time_bnds	(time, bnds)	object	dask.array<chunksize=(1, 2), meta=np.ndarray>	lev_bnds	(time, lev, bnds)	float64	dask.array<chunksize=(240, 35, 2), meta=np.ndarray>
lat_bnds	(time, lat, bnds)	float64	dask.array<chunksize=(240, 180, 2), meta=np.ndarray>																				
lon_bnds	(time, lon, bnds)	float64	dask.array<chunksize=(240, 360, 2), meta=np.ndarray>																				
thetao	(time, lev, lat, lon)	float32	dask.array<chunksize=(1, 35, 180, 360), meta=np.ndarray>																				
time_bnds	(time, bnds)	object	dask.array<chunksize=(1, 2), meta=np.ndarray>																				
lev_bnds	(time, lev, bnds)	float64	dask.array<chunksize=(240, 35, 2), meta=np.ndarray>																				
▼ Attributes:																							
<table><tr><td>title :</td><td colspan="3">NOAA GFDL GFDL-CM4 model output prepared for CMIP6 OMIP experiment forced by CORE Inter-Annual Forcing (CIAF) version 2.0 and initialized with observed physical and biogeochemical ocean data</td></tr><tr><td>history :</td><td colspan="3">File was processed by fremetar (GFDL analog of CMOR). TripleID: [exper_id_TFN7SZc1yU,realiz_id_HX30DNwWjL,run_id_aOuCxFhBdP]</td></tr></table>				title :	NOAA GFDL GFDL-CM4 model output prepared for CMIP6 OMIP experiment forced by CORE Inter-Annual Forcing (CIAF) version 2.0 and initialized with observed physical and biogeochemical ocean data			history :	File was processed by fremetar (GFDL analog of CMOR). TripleID: [exper_id_TFN7SZc1yU,realiz_id_HX30DNwWjL,run_id_aOuCxFhBdP]														
title :	NOAA GFDL GFDL-CM4 model output prepared for CMIP6 OMIP experiment forced by CORE Inter-Annual Forcing (CIAF) version 2.0 and initialized with observed physical and biogeochemical ocean data																						
history :	File was processed by fremetar (GFDL analog of CMOR). TripleID: [exper_id_TFN7SZc1yU,realiz_id_HX30DNwWjL,run_id_aOuCxFhBdP]																						

Data exploration



Encourage the use of Data/Metadata **cataloging capability**.

E.g. intake-esm: a data cataloging utility from Pangeo that uses

ESM collection specifications file and is built on top of [intake](#), [pandas](#), and [xarray](#)

GFDL's CatalogBuilder- A python API to build custom data inventory catalogs (e.g. intake-esm) from data hosted in the cloud or your local file system.



Ack: A.Radhakrishnan, Rebecca Monge



A.Radhakrishnan et al, Building blocks for exascale computing
at GFDL, Presented at 6th ENES HPC Workshop 2020



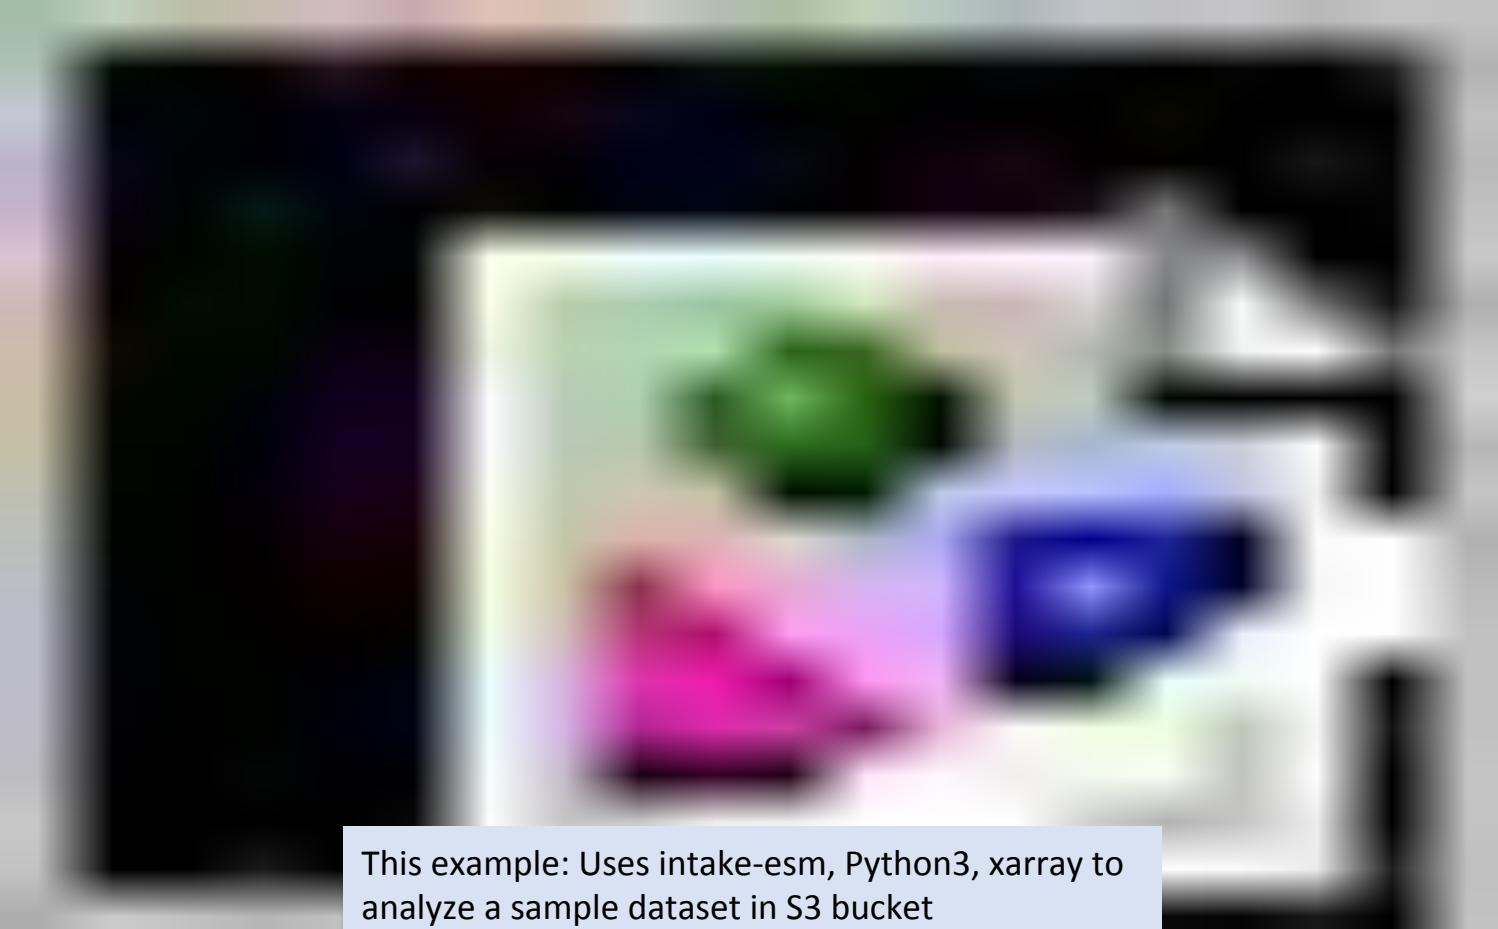
Traceability

Useful resources to access S3 using Python

[Our S3 data inventory catalog](#) (updated regularly)

[Example](#) script that uses the above data catalogs (intake-esm in this example) to fit into our analysis seamlessly.

Example notebook to access S3 using Python



This example: Uses intake-esm, Python3, xarray to analyze a sample dataset in S3 bucket

Tools to monitor resources and applications



Amazon Cloudwatch

To monitor services and provide a unified view of the operational health of the AWS resources

CloudWatch: **EC2** ▾

jhub-instances

ALARM 0 INSUFFICIENT DATA 0

CPU Utilization Ave

Percent

57.63

30.3

3

18:00 12-05 19:05 20:00

2020-12-05 19:05 UTC

i-0a1f2
i-049f
i-048c

1. i-04f9f4af93b9865de 57.6
2. i-048c2507eebf137d6 9

CloudWatch

Dashboards

Alarms

ALARM 1

INSUFFICIENT 0

OK 1

Billing

Logs

Log groups

Insights

Metrics

Explorer **NEW**

Events

Rules

Event Buses

ServiceLens

Service Map

Traces



serverless compute service: lets you run code without provisioning servers

/aws/lambda/lambdatest1

▼ Log group details

Retention	Creation time	Stored bytes
Never expire	3 days ago	33.46 KB
KMS key ID	Metric filters	Subscription filters
-	0	0

Log streams Metric filters Subscription filters Contributor insights

Log streams (57)

Filter log streams or try prefix search

Log stream	Last event time
2020/12/05/[\$LATEST]0541ed21ff804cde9bc23fd5bfae619b	2020-12-05 15:25:24 (UTC-05:00)
2020/12/02/[\$LATEST]cf1553e21d1e42eb2c686020e9c1b7e	2020-12-02 16:20:28 (UTC-05:00)

AWS Cost Management

Home

Cost Explorer

Reports

Budgets

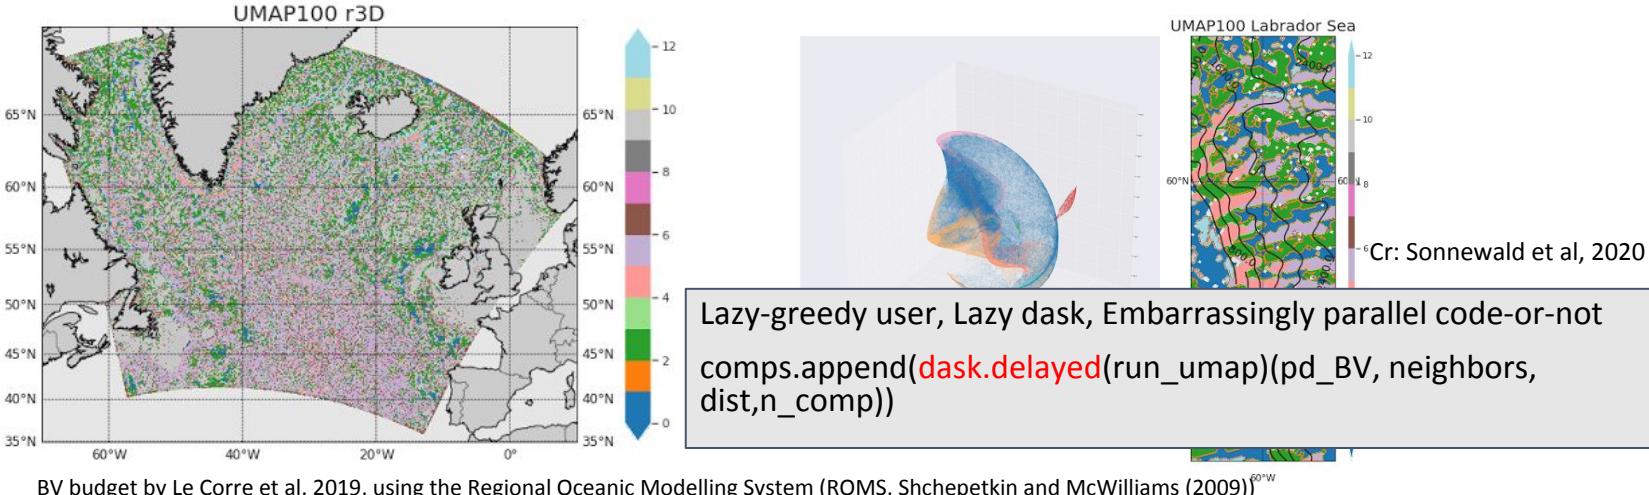
Create a budget



- Step 1 Select budget type
- Step 2 Set your budget
- Step 3 Configure thresholds
- Step 4 Confirm budget

Ongoing/Future work

- Cost optimizations and instance configurations. e.g requester pays cost model.
- Further testing and performance tuning
- Robust strategies for configuring Dask array chunking
- Explore use of cloud based analysis in more challenging ML applications
- Complete deployment of Earth System Grid Federation software stack in AWS/EKS.
- Encourage users to use official CMIP6 data citation from further_info_url global attribute of NetCDF data in S3.



Climate in the cloud, V.Balaji, A.Radhakrishnan



Thank you!

Many thanks to GFDL, ESGF, Pangeo community and our AWS colleagues and collaborators.

Additional references

- <https://docs.aws.amazon.com/AmazonS3/latest/dev/example-walkthroughs-managing-access-example2.html>
- <https://jupyter.org/>
- <https://docs.aws.amazon.com/eks/latest/userguide/getting-started.html>
- <https://zero-to-jupyterhub.readthedocs.io/en/latest/kubernetes/amazon/step-zero-aws-eks.html>
- <https://docs.aws.amazon.com/>
- <https://github.com/aradhakrishnanGFDL/gfdl-aws-analysis/tree/community/examples>